

dplyr Practical Solutions

Jumping Rivers

We'll start by loading the necessary packages and data sets

```
library("dplyr")
library("ggplot2")
data(okcupid, package = "jrTidyverse")
```

Summarising the data

In this section, we will gradually chain the commands together. We'll start things off, by calculating the average income

```
new_data = okcupid %>%
  summarise(ave_income = mean(income))
new_data
```

```
## # A tibble: 1 x 1
##   ave_income
##   <dbl>
## 1    104395.
```

1. Alter the above command to calculate the median income (as well as the mean).

```
okcupid %>%
  summarise(ave_income = mean(income),
            med_income = median(income))
```

```
## # A tibble: 1 x 2
##   ave_income med_income
##   <dbl>         <dbl>
## 1    104395.         50000
```

2. Use the `group_by()` to calculate the average incomes conditional on the answer to the `pets` question.

```
okcupid %>%
  group_by(pets) %>%
  summarise(ave_income = mean(income))
```

```
## # A tibble: 16 x 2
##   pets                ave_income
##   <chr>                <dbl>
## 1 dislikes cats      159500
## 2 dislikes dogs      66154.
## 3 dislikes dogs and dislikes cats 176154.
## 4 dislikes dogs and has cats      93953.
## 5 dislikes dogs and likes cats    103956.
## 6 has cats           84498.
## 7 has dogs          112540.
## 8 has dogs and dislikes cats    104895.
## 9 has dogs and has cats      97995.
## 10 has dogs and likes cats     87432.
## 11 likes cats         69234.
```

```
## 12 likes dogs 119483.
## 13 likes dogs and dislikes cats 99667.
## 14 likes dogs and has cats 90905.
## 15 likes dogs and likes cats 106814.
## 16 <NA> 106222.
```

4. The `arrange()` function is used to sort a tibble, .e.g

```
... %>%
  arrange(ave_income)
```

will arrange the tibble from smallest to largest. Arrange the tibble from **largest** to smallest in terms of average income.

```
(df = okcupid %>%
  group_by(pets) %>%
  summarise(ave_income = mean(income)) %>%
  arrange(desc(ave_income))
)
```

```
## # A tibble: 16 x 2
##   pets ave_income
##   <chr>      <dbl>
## 1 dislikes dogs and dislikes cats 176154.
## 2 dislikes cats 159500
## 3 likes dogs 119483.
## 4 has dogs 112540.
## 5 likes dogs and likes cats 106814.
## 6 <NA> 106222.
## 7 has dogs and dislikes cats 104895.
## 8 dislikes dogs and likes cats 103956.
## 9 likes dogs and dislikes cats 99667.
## 10 has dogs and has cats 97995.
## 11 dislikes dogs and has cats 93953.
## 12 likes dogs and has cats 90905.
## 13 has dogs and likes cats 87432.
## 14 has cats 84498.
## 15 likes cats 69234.
## 16 dislikes dogs 66154.
```

Creating columns with `mutate()`

1. The `floor()` function rounds down to the nearest integer. To round to the nearest 10, we use the trick

```
floor(61/10)*10
```

```
## [1] 60
```

```
floor(119/10)*10
```

```
## [1] 110
```

We can use the `mutate()` function to create a new column that contains the persons age (to the decade), i.e. 50, 60, 70, etc. The `mutate()` function isn't directly in the notes, but it is relatively easy to understand. It creates a new column with the given name, based on manipulation of existing columns. So we could create this new column **decade**.

```
okcupid %>%
  mutate(decade = floor(age/10) * 10)
```

```
## # A tibble: 11,504 x 22
##   age body_type diet drinks drugs education ethnicity height income
##   <int> <chr>    <chr> <chr> <chr> <chr>    <chr>    <int> <int>
## 1 35 average most~ often some~ working ~ white      70 8.00e4
## 2 23 thin vege~ socia~ <NA> working ~ white      71 2.00e4
## 3 28 average most~ socia~ never graduate~ white      72 4.00e4
## 4 30 skinny most~ socia~ never graduate~ white      66 3.00e4
## 5 29 thin most~ socia~ never working ~ hispanic~ 62 5.00e4
## 6 40 fit <NA> socia~ <NA> graduate~ white      71 6.00e4
## 7 31 thin stri~ socia~ some~ dropped ~ <NA>      67 1.00e6
## 8 22 athletic most~ rarely never working ~ asian      65 2.00e4
## 9 35 athletic most~ socia~ some~ graduate~ native a~ 73 1.50e5
## 10 31 curvy most~ socia~ never graduate~ indian      61 5.00e4
## # ... with 11,494 more rows, and 13 more variables: job <chr>,
## # last_online <dtm>, location <chr>, offspring <chr>,
## # orientation <chr>, pets <chr>, religion <chr>, sex <chr>, sign <chr>,
## # smokes <chr>, speaks <chr>, status <chr>, decade <dbl>
```

2. Since this data set has high earners, use `filter()` to remove the top 5% of earners. Hint: `quantile(income, probs = 0.95)` will give you the 95%-tile of income

```
okcupid %>%
  mutate(decade = floor(age/10)*10) %>%
  filter(income < quantile(income, probs = 0.95))
```

```
## # A tibble: 10,786 x 22
##   age body_type diet drinks drugs education ethnicity height income
##   <int> <chr>    <chr> <chr> <chr> <chr>    <chr>    <int> <int>
## 1 35 average most~ often some~ working ~ white      70 80000
## 2 23 thin vege~ socia~ <NA> working ~ white      71 20000
## 3 28 average most~ socia~ never graduate~ white      72 40000
## 4 30 skinny most~ socia~ never graduate~ white      66 30000
## 5 29 thin most~ socia~ never working ~ hispanic~ 62 50000
## 6 40 fit <NA> socia~ <NA> graduate~ white      71 60000
## 7 22 athletic most~ rarely never working ~ asian      65 20000
## 8 35 athletic most~ socia~ some~ graduate~ native a~ 73 150000
## 9 31 curvy most~ socia~ never graduate~ indian      61 50000
## 10 21 fit most~ rarely never working ~ white      71 20000
## # ... with 10,776 more rows, and 13 more variables: job <chr>,
## # last_online <dtm>, location <chr>, offspring <chr>,
## # orientation <chr>, pets <chr>, religion <chr>, sex <chr>, sign <chr>,
## # smokes <chr>, speaks <chr>, status <chr>, decade <dbl>
```

3. To help with plotting, convert the decade column into a character using the `as.character()` function. This can be achieved via `mutate(decade = as.character(decade))`

```
(df = okcupid %>%
  mutate(decade = floor(age/10)*10) %>%
  filter(income < quantile(income, probs = 0.95)) %>%
  mutate(decade = as.character(decade))
)
```

```
## # A tibble: 10,786 x 22
```

```
##      age body_type diet  drinks drugs education ethnicity height income
##      <int> <chr>      <chr> <chr> <chr> <chr>      <chr>      <int> <int>
##  1     35 average  most~ often  some~ working ~ white        70 80000
##  2     23 thin    vege~ socia~ <NA> working ~ white        71 20000
##  3     28 average  most~ socia~ never graduate~ white        72 40000
##  4     30 skinny  most~ socia~ never graduate~ white        66 30000
##  5     29 thin    most~ socia~ never working ~ hispanic~    62 50000
##  6     40 fit      <NA> socia~ <NA> graduate~ white        71 60000
##  7     22 athletic most~ rarely never working ~ asian        65 20000
##  8     35 athletic most~ socia~ some~ graduate~ native a~    73 150000
##  9     31 curvy   most~ socia~ never graduate~ indian        61 50000
## 10     21 fit     most~ rarely never working ~ white        71 20000
## # ... with 10,776 more rows, and 13 more variables: job <chr>,
## #   last_online <dtm>, location <chr>, offspring <chr>,
## #   orientation <chr>, pets <chr>, religion <chr>, sex <chr>, sign <chr>,
## #   smokes <chr>, speaks <chr>, status <chr>, decade <chr>
```